Programming Assignment 1: Parallel Computing Odd-Even Sort and Pearson Correlation Coefficient

Course: 4DT906 - Parallel Computing

February 3, 2026

Abstract

This report presents parallel implementations of two algorithms using MPI (Message Passing Interface): the odd-even transposition sort and Pearson correlation coefficient calculation. Both implementations demonstrate significant speedup compared to their sequential counterparts. The odd-even sort achieves correct sorting through parallel compare-exchange operations, while the parallel Pearson correlation efficiently distributes the computation of correlation pairs across multiple processes. Performance benchmarks show speedup factors of up to 3.56x with 4 processes.

Contents

1 Introduction

Parallel computing has become essential for handling large-scale computational problems efficiently. This assignment explores two fundamental parallel algorithms:

1. **Odd-Even Transposition Sort**: A comparison-based sorting algorithm particularly suited for parallel execution due to its regular communication pattern.

2. **Pearson Correlation Coefficient (PCC)**: A statistical measure of linear correlation between variables, computationally intensive for large datasets but highly parallelizable.

Both algorithms were implemented using MPI, which provides a standard message-passing interface for distributed-memory parallel computing.

2 Odd-Even Transposition Sort

2.1 Algorithm Description

The odd-even transposition sort is a variation of bubble sort that alternates between two phases:

- **Odd phase**: Compares and exchanges elements at odd-indexed positions with their right neighbors

- **Even phase**: Compares and exchanges elements at even-indexed positions with their right neighbors

In the parallel version, each process holds a subset of the data, and processes perform compare-exchange operations with their neighbors.

2.2 Sequential Implementation

The sequential implementation performs n iterations (where n is the number of elements), alternating between odd and even phases. In each phase, adjacent elements are compared and swapped if needed:

```
for (int i = 1; i <= s; i++) {
    for (int j = i % 2; j < s-1; j = j + 2) {
        if (numbers[j] > numbers[j + 1]) {
            std::swap(numbers[j], numbers[j + 1]);
        }
    }
}
```

Listing 1: Sequential Odd-Even Sort Core Loop

Time Complexity: $O(n^2)$ in the worst case, where n is the number of elements.

2.3 Parallel Implementation

The parallel implementation divides the data equally among p processes. Each process:

1. Receives its local data portion via `MPI_Scatter`

2. Sorts its local data using standard sort

3. Performs p iterations of odd-even phases

4. In each phase, exchanges data with neighboring processes and keeps appropriate values

5. Gathers results back to root using `MPI_Gather`

The key operations are `compare_exchange_low` and `compare_exchange_high`:

- **compare_exchange_low**: Process keeps the smaller half of merged data

- **compare_exchange_high**: Process keeps the larger half of merged data

```cpp
void compare_exchange_low(std::vector<int>& local_numbers,
                          int partner, int n_local) {
    std::vector<int> recv_numbers(n_local);

    MPI_Sendrecv(local_numbers.data(), n_local, MPI_INT, partner, 0,
                 recv_numbers.data(), n_local, MPI_INT, partner, 0,
                 MPI_COMM_WORLD, MPI_STATUS_IGNORE);

    // Merge and keep the smaller n_local elements
    std::vector<int> merged;
    merged.insert(merged.end(), local_numbers.begin(),
                  local_numbers.end());
    merged.insert(merged.end(), recv_numbers.begin(),
                  recv_numbers.end());
    std::sort(merged.begin(), merged.end());

    // Keep the lower half
    for (int i = 0; i < n_local; i++) {
        local_numbers[i] = merged[i];
    }
}
```

Listing 2: Parallel Compare-Exchange Operation

2.4 Performance Analysis

For n elements and p processes:

- Each process handles n/p elements

- Local sort: $O(\frac{n}{p} \log \frac{n}{p})$

- Parallel phases: $O(p \cdot \frac{n}{p} \log \frac{n}{p})$

- Communication: $O(p \cdot \frac{n}{p})$ per iteration

Expected Speedup: Near-linear for moderate p values, with communication overhead becoming significant as p increases.

3 Pearson Correlation Coefficient

3.1 Algorithm Description

The Pearson correlation coefficient measures linear correlation between pairs of variables. For a matrix with m rows and n columns, we compute correlations between all pairs of rows, resulting in $\frac{m(m-1)}{2}$ correlation values.

The computation involves three main steps:

1. Calculate row means

2. Compute mean-adjusted values and standard deviations

3. Calculate pairwise correlations using:

$$r_{ij} = \frac{\sum_{k=1}^{n}(x_{ik} - \bar{x}_i)(x_{jk} - \bar{x}_j)}{\sigma_i \sigma_j}$$

3.2 Sequential Implementation

The sequential version processes all correlation pairs in nested loops:

```
for(sample1 = 0; sample1 < ROWS-1; sample1++){
    for(sample2 = sample1+1; sample2 < ROWS; sample2++){
        sum = 0.0;
        for(i = 0; i < COLS; i++){
            sum += mm[sample1 * COLS + i] * mm[sample2 * COLS + i];
        }
        r = sum / (std[sample1] * std[sample2]);
        output[index] = r;
    }
}
```

Listing 3: Sequential Pearson Correlation

Time Complexity: $O(m^2 n)$ where m is the number of rows and n is the number of columns.

3.3 Parallel Implementation

The parallel implementation distributes the correlation pairs among processes:

1. All processes compute row means and standard deviations (small overhead, simplifies implementation)

2. Work is divided: each process computes a subset of the $\frac{m(m-1)}{2}$ correlation pairs

3. Results are gathered using `MPI_Gatherv` to handle potentially uneven work distribution

```
// Divide work among processes
int local_size = cor_size / size;
int remainder = cor_size % size;
int local_start = rank * local_size +
                  (rank < remainder ? rank : remainder);
int local_count = local_size + (rank < remainder ? 1 : 0);
int local_end = local_start + local_count;
```

Listing 4: Parallel Work Distribution

3.4 Performance Analysis

For an $m \times n$ matrix with p processes:

- Preprocessing (mean/std): $O(mn)$ (all processes)

- Correlation computation: $O(\frac{m^2 n}{p})$ per process

- Communication: $O(\frac{m^2}{p})$ for gathering results

Expected Speedup: Near-linear for large matrices, as computation dominates communication overhead.

4 Experimental Results

4.1 Pearson Correlation Coefficient Benchmark

The benchmark was run on matrices of increasing sizes, comparing sequential and parallel (4 processes) implementations:

Matrix Size	Sequential (s)	Parallel (s)	Speedup
64 × 64	0.0006	0.0003	2.00×
128 × 128	0.0042	0.0015	2.80×
256 × 256	0.0324	0.0103	3.15×
512 × 512	0.2554	0.0755	3.38×
1024 × 1024	2.0209	0.5713	3.54×
2048 × 2048	16.1488	4.6032	3.51×
4096 × 4096	128.9257	36.2020	3.56×

Table 1: PCC Performance Results (4 Processes)

4.2 Analysis

- **Speedup Trend**: Speedup increases from 2.00× to 3.56× as problem size grows, approaching the theoretical maximum of 4× for 4 processes.

- **Communication Overhead**: More pronounced in smaller problems (64×64), where speedup is only 2.00×.

- **Scalability**: For large problems (4096×4096), the parallel version achieves 3.56× speedup, demonstrating good scalability.

- **Efficiency**: Parallel efficiency ranges from 50% (small problems) to 89% (large problems), calculated as $\frac{\text{Speedup}}{p}$.

4.3 Odd-Even Sort Results

The parallel odd-even sort successfully sorts 100,000 elements using 4 processes. Testing confirmed:

- **Correctness**: Output is correctly sorted (verified with `std::is_sorted`)

- **Performance**: Execution time of approximately 0.115 seconds for 100,000 elements

5 Discussion

5.1 Advantages of Parallel Implementations

1. **Significant Speedup**: Both implementations achieve substantial performance improvements, especially for large datasets.

2. **Scalability**: The PCC implementation shows good scalability, with efficiency increasing as problem size grows.

3. **Correctness**: Validation confirms that parallel versions produce identical results to sequential versions.

5.2   Limitations and Challenges

1. **Communication Overhead**: For small problems, communication costs dominate, limiting speedup.

2. **Memory Requirements**: Odd-even sort's compare-exchange operations require temporary buffers, doubling memory usage during exchanges.

3. **Load Balancing**: In PCC, remainder work distribution may cause slight load imbalance.

5.3   Potential Optimizations

1. **Hybrid Approach**: Use OpenMP for shared-memory parallelism within nodes and MPI between nodes.

2. **Asynchronous Communication**: Overlap computation with communication using non-blocking MPI operations.

3. **Better Data Distribution**: For PCC, distribute based on row-pair indices to improve cache locality.

4. **Reduced Redundancy**: In PCC, only root process needs to compute mean/std, then broadcast results.

6   Conclusion

This assignment successfully demonstrated parallel implementations of odd-even sort and Pearson correlation coefficient using MPI. Key achievements include:

- **Correct Implementations**: Both algorithms produce accurate results verified against sequential versions.

- **Performance Gains**: PCC achieves up to $3.56\times$ speedup with 4 processes, with efficiency improving for larger problems.

- **Understanding**: Gained practical experience with MPI collective operations, data distribution, and parallel algorithm design.

The results confirm that parallelization can significantly reduce execution time for computationally intensive tasks, particularly when the problem size is large enough to amortize communication overhead. The implementations serve as a foundation for understanding more complex parallel algorithms and optimization techniques.

References

- MPI Forum. MPI: A Message-Passing Interface Standard, Version 3.1. 2015.

- Quinn, M. J. Parallel Programming in C with MPI and OpenMP. McGraw-Hill, 2004.

- Course Material: PA1_26.pdf - Programming Assignment 1 Instructions